



# [Hands-on] Fast Training ImageNet on on-demand EC2 GPU instances with Horovod

Author: Daekeun Kim (daekeun@amazon.com)

## Goal

This document is for people who need distributed GPU training using Horovod for experimental purposes. Many steps are similar to what mentioned in Julien Simon's article(<https://medium.com/@julsimon/imagenet-part-1-going-on-an-adventure-c0a62976dc72>) and AWS

Documentation(<https://docs.aws.amazon.com/dlami/latest/devguide/tutorial-horovod-tensorflow.html>). So I recommend you to view these articles first. If there are some things that aren't going well (e.g., Downloading the dataset does not work, How to convert the raw data to the TFRecord feature set?, How to fix the error `ModuleNotFoundError: No module named 'cv2'` ) please refer this document.

## Introduction

For data preparation and data transformation, we do not need to use a GPU instance such as p2 and p3. Instead, we can start much cheaper instances like `t2.large` instance with 1.0TB EBS volume.

For distributed training, we need to use multiple GPU instances like p2, p3, g3 and g4.

You can skip step 1 if you do not want to invent the wheel again because I have stored everything in my s3 bucket.

- `s3://dataset-image/imagenet/raw` (raw jpeg)
- `s3://dataset-image/imagenet/tfrecord` (TFRecord before resizing)
- `s3://dataset-image/imagenet/tfrecord-resized` (TFRecord after resizing to 224x224)
- `s3://dataset-image/imagenet/recordio` (RecordIO after resizing to 256x256)
  - The reason I did not resize 224x224 is that the below article shows different validation accuracy for resizing strategy.
  - <https://forums.fast.ai/t/impact-of-image-resizing-on-model-training-time-and-performance/1980>

Please let me know if you want to access the bucket because I did not grant any public access.

## Step 1. Downloading and Transformation

### Setting up an EC2 instance for Data Transformation

- Create an EC2 instance for storing ImageNet dataset (Ubuntu 18.04 or 16.04. Linux is also available). `t2.micro` is also available, but `t2.large` is recommended due to memory size. Note that we do not need large storage size since we will make another EBS volume to attach the EC2 instance.
- Create an EBS volume (1.0TB) for ImageNet dataset and then attach the volume it to your EC2 instance. ImageNet consists of 138GB for training set and 6.3GB for validation set, but we need an additional space since we need to extract tar files as well as need to transform it to the feature sets like TFRecord and RecordIO. Here is an example command using AWS CLI.

```
$ aws ec2 create-volume \ --size 1000 \ --region
[YOUR_AWS_REGION] \ --availability-zone [YOUR_AZ] \ --volume-
type sc1 \ --tag-specifications 'ResourceType=volume,Tags=
[{{Key=Name,Value=ImageNet}}]' $ aws ec2 attach-volume \ --volume-
id vol-[YOUR_EC2_volume_id] \ --instance-id i-
[YOUR_EC2_instance_id] \ --device /dev/sdf
```

- You can use the TensorFlow's download script([https://github.com/tensorflow/models/blob/master/research/inception/inception/data/download\\_imagenet.sh](https://github.com/tensorflow/models/blob/master/research/inception/inception/data/download_imagenet.sh)) by exporting your username and access key.

```
$ export IMAGENET_USERNAME=[YOUR_USERNAME] $ export
IMAGENET_ACCESS_KEY=[YOUR_ACCESS_KEY] $ cd imagenet/data $ mv
imagenet_2012_validation_synset_labels.txt synsets.txt $ nohup
bash download_imagenet.sh . synsets.txt >& download.log &
```

## Method 2 (Alternative method if Method 1 does not work)

- Download ImageNet dataset manually.

```
$ nohup wget http://www.image-
net.org/challenges/LSVRC/2012/nnoupb/ILSVRC2012_img_train.tar & $
nohup wget http://www.image-
net.org/challenges/LSVRC/2012/nnoupb/ILSVRC2012_bbox_train_v2.tar.
& $ nohup wget http://www.image-
net.org/challenges/LSVRC/2012/nnoupb/ILSVRC2012_img_val.tar & $
nohup wget http://www.image-
net.org/challenges/LSVRC/2012/nnoupb/ILSVRC2012_bbox_val_v3.tgz &
```

- Extract the validation set

```
$ mkdir validation $ mv ILSVRC2012_img_val.tar validation $ cd
validation $ tar xf ILSVRC2012_img_val.tar
```

- After extracting the validation set, move jpeg files(ILSVRC2012\_val\_000000001.JPEG, ..., ILSVRC2012\_val\_00050000.JPEG) in 1,000 directories using the following script;  
[https://github.com/juliensimon/aws/blob/master/mxnet/imagenet/build\\_validation\\_tree.sh](https://github.com/juliensimon/aws/blob/master/mxnet/imagenet/build_validation_tree.sh) (Each directory means the unique category like ).

n01728572

- Please refer to [https://github.com/aws-samples/deep-learning-models/blob/master/utils/tensorflow/preprocess\\_imagenet.py](https://github.com/aws-samples/deep-learning-models/blob/master/utils/tensorflow/preprocess_imagenet.py) (code) and [https://docs.aws.amazon.com/ko\\_kr/dlami/latest/devguide/tutorial-horovod-tensorflow.html](https://docs.aws.amazon.com/ko_kr/dlami/latest/devguide/tutorial-horovod-tensorflow.html) (document).
  - `python preprocess_imagenet.py \ --local_scratch_dir=[YOUR DIRECTORY] \ --imagenet_username=[imagenet account] \ --imagenet_access_key=[imagenet access key]`
  - `python tensorflow_image_resizer.py \ -d imagenet \ -i [PATH TO TFRECORD TRAINING DATASET] \ -o [PATH TO RESIZED TFRECORD TRAINING DATASET] \ --subset_name train \ --num_preprocess_threads 60 \ --num_intra_threads 2 \ --num_inter_threads 2`
- [Additional Notes] The original document uses the small number of intra-op(multiple threads within one op; for example, while doing matrix multiplication operation we can divide the op by multiple threads) and inter-op(thread-pool size per one executor) such that `--num_intra_threads 2 \ --num_inter_threads 2`. But, you can give higher number of intra-op and inter-op.

## Backing up and Copying to S3

- After data transformation, create a new bucket and sync or copy feature sets to the bucket.
- Create a snapshot of the EBS volume.

## Step 2. Training ResNet-50 Model with Horovod

[Before get started] If you just want to train on a single machine, you may refer to <https://medium.com/@julsimon/imagenet-part-2-the-road-goes-ever-on-and-on-578f09a749f9> (RecordIO) and <https://github.com/tensorflow/models/tree/master/official/r1/resnet> (TFRecord)

- hvd\_train\_log (64 GPUS; 8 p3dn.24xlarge instances)

```
- Step Epoch Speed Loss FinLoss LR - 0 0.0 1907.3 6.920 8.259  
0.00100 - 1 0.0 5164.9 6.935 8.274 0.00920 - 50 0.6 43926.5  
6.206 7.522 0.41119 - ... - 6950 88.9 43552.2 0.783 1.185  
0.00125 - 7000 89.5 41958.4 0.624 1.027 0.00023 - Finished in  
2685.1825189590454
```

- eval\_hvd\_train.log (64 GPUS; 8 p3dn.24xlarge instances)